

MovieRecom

April 9, 2025

```
[3]: # Step 1: Importing Essential Libraries and Loading Data
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics.pairwise import cosine_similarity
from scipy.sparse import csr_matrix
```

```
[4]: # Load movie and rating data (ensure CSV files are in your working directory)
movie_data = pd.read_csv("movies.csv")
rating_data = pd.read_csv("ratings.csv")
```

```
[5]: # Display basic information about the movies dataset
print("Movies Data Info:")
print(movie_data.info())
print(movie_data.describe())
print(movie_data.head())

print("\nRatings Data Info:")
print(rating_data.info())
print(rating_data.head())
```

Movies Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10329 entries, 0 to 10328

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	movieId	10329 non-null	int64
1	title	10329 non-null	object
2	genres	10329 non-null	object

dtypes: int64(1), object(2)

memory usage: 242.2+ KB

None

	movieId
count	10329.000000
mean	31924.282893
std	37734.741149
min	1.000000

```

25%      3240.000000
50%      7088.000000
75%     59900.000000
max     149532.000000

movieId      title \
0           1      Toy Story (1995)
1           2      Jumanji (1995)
2           3      Grumpier Old Men (1995)
3           4      Waiting to Exhale (1995)
4           5      Father of the Bride Part II (1995)

genres
0  Adventure|Animation|Children|Comedy|Fantasy
1      Adventure|Children|Fantasy
2      Comedy|Romance
3      Comedy|Drama|Romance
4      Comedy

```

Ratings Data Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 105339 entries, 0 to 105338

Data columns (total 4 columns):

```

#   Column      Non-Null Count  Dtype
---  -----  -
0   userId    105339 non-null    int64
1   movieId    105339 non-null    int64
2   rating     105339 non-null    float64
3   timestamp  105339 non-null    int64

```

dtypes: float64(1), int64(3)

memory usage: 3.2 MB

None

```

userId  movieId  rating  timestamp
0        1       16     4.0  1217897793
1        1       24     1.5  1217895807
2        1       32     4.0  1217896246
3        1       47     4.0  1217896556
4        1       50     4.0  1217896523

```

```

[6]: # Step 2: Data Pre-processing - One-Hot Encoding of Genres
# Split the 'genres' column (which contains strings with genres separated by
↳ '/'')
movie_data['genres_list'] = movie_data['genres'].apply(lambda x: x.split('|'))

# Define the list of genres (same as used in the R code)
all_genres = ["Action", "Adventure", "Animation", "Children", "Comedy", "Crime",
              "Documentary", "Drama", "Fantasy", "Film-Noir", "Horror",
↳ "Musical",

```

```

        "Mystery", "Romance", "Sci-Fi", "Thriller", "War", "Western"]

# Create one-hot encoding for each genre
for genre in all_genres:
    movie_data[genre] = movie_data['genres_list'].apply(lambda genres: 1 if
        genre in genres else 0)

# Create a search matrix that includes movieId, title, and genre columns
search_matrix = movie_data[['movieId', 'title'] + all_genres]
print("\nSearch Matrix (first 5 rows):")
print(search_matrix.head())

# Step 3: Construct the Rating Matrix
# Create a pivot table where rows are users and columns are movies
rating_matrix = rating_data.pivot(index='userId', columns='movieId',
    values='rating')
print("\nRating Matrix shape:", rating_matrix.shape)

# (Optional) Convert to a sparse matrix if needed for efficiency:
rating_matrix_sparse = csr_matrix(rating_matrix.fillna(0).values)

# Step 4: Exploring Similarity - Users and Movies
# For users: compute cosine similarity on the first 4 users (fill NaNs with 0)
user_similarity = cosine_similarity(rating_matrix.iloc[:4].fillna(0))
print("\nUser Similarity (first 4 users):")
print(user_similarity)

```

Search Matrix (first 5 rows):

	movieId		title	Action	Adventure	Animation	\
0	1		Toy Story (1995)	0	1	1	
1	2		Jumanji (1995)	0	1	0	
2	3		Grumpier Old Men (1995)	0	0	0	
3	4		Waiting to Exhale (1995)	0	0	0	
4	5		Father of the Bride Part II (1995)	0	0	0	

	Children	Comedy	Crime	Documentary	Drama	Fantasy	Film-Noir	Horror	\
0	1	1	0	0	0	1	0	0	
1	1	0	0	0	0	1	0	0	
2	0	1	0	0	0	0	0	0	
3	0	1	0	0	1	0	0	0	
4	0	1	0	0	0	0	0	0	

	Musical	Mystery	Romance	Sci-Fi	Thriller	War	Western
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	1	0	0	0	0

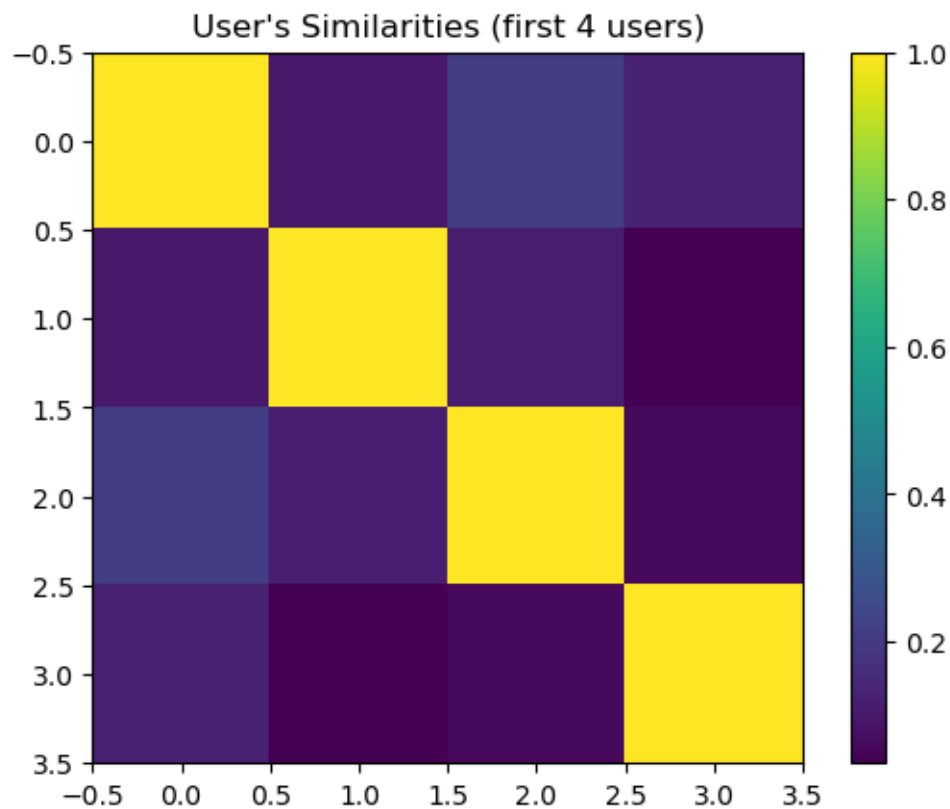
3	0	0	1	0	0	0	0
4	0	0	0	0	0	0	0

Rating Matrix shape: (668, 10325)

User Similarity (first 4 users):

```
[[1.          0.10111327 0.21004361 0.12876575]
 [0.10111327 1.          0.11555911 0.0346102 ]
 [0.21004361 0.11555911 1.          0.05820771]
 [0.12876575 0.0346102  0.05820771 1.          ]]
```

```
[7]: plt.figure()
plt.imshow(user_similarity, interpolation='nearest', cmap='viridis')
plt.title("User's Similarities (first 4 users)")
plt.colorbar()
plt.show()
```



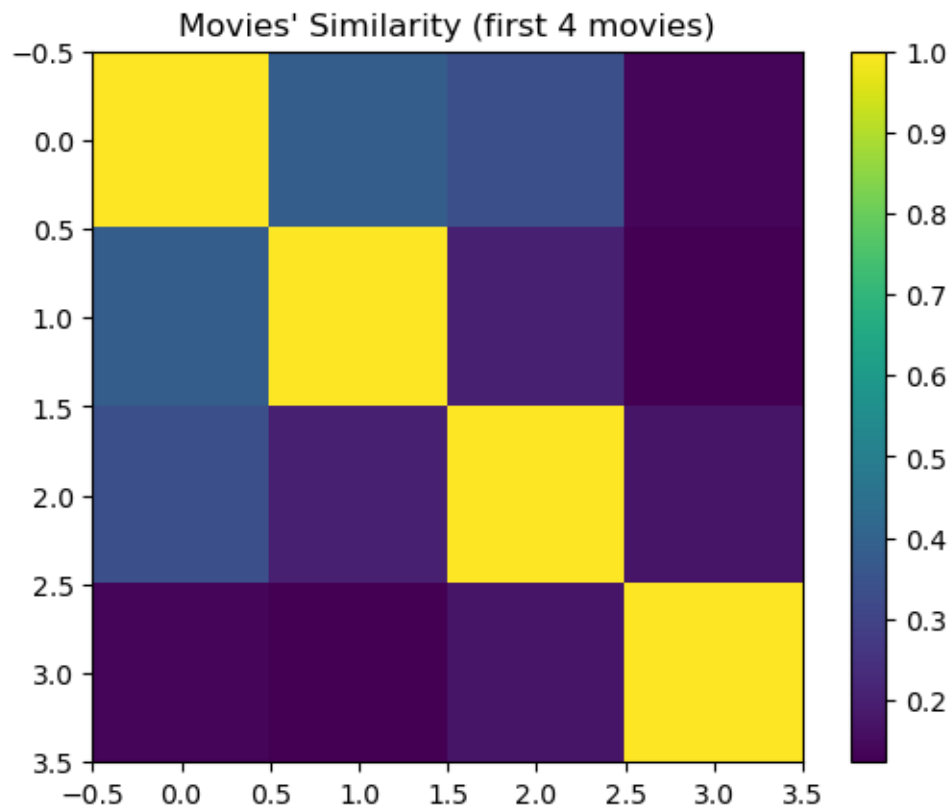
```
[8]: # For movies: compute cosine similarity on the first 4 movies
# (Transposing so that movies are treated as vectors)
movie_similarity = cosine_similarity(rating_matrix.iloc[:, :4].fillna(0).T)
print("\nMovie Similarity (first 4 movies):")
```

```
print(movie_similarity)
```

Movie Similarity (first 4 movies):

```
[[1.          0.38306843 0.33745279 0.13472433]
 [0.38306843 1.          0.19920682 0.12337653]
 [0.33745279 0.19920682 1.          0.17336631]
 [0.13472433 0.12337653 0.17336631 1.          ]]
```

```
[9]: plt.figure()
plt.imshow(movie_similarity, interpolation='nearest', cmap='viridis')
plt.title("Movies' Similarity (first 4 movies)")
plt.colorbar()
plt.show()
```



```
[11]: # Step 5: Unique Ratings and Table of Ratings
# Get unique rating values (ignoring NaNs)
unique_ratings = np.unique(rating_matrix.values[~np.isnan(rating_matrix.
    ↪ values)])
print("\nUnique Ratings:", unique_ratings)
```

```
# Count the occurrence of each rating from the raw rating data
table_of_ratings = rating_data['rating'].value_counts().sort_index()
print("\nTable of Ratings:")
print(table_of_ratings)
```

Unique Ratings: [0.5 1. 1.5 2. 2.5 3. 3.5 4. 4.5 5.]

Table of Ratings:

```
rating
0.5    1198
1.0    3258
1.5    1567
2.0    7943
2.5    5484
3.0   21729
3.5   12237
4.0   28880
4.5    8187
5.0   14856
Name: count, dtype: int64
```

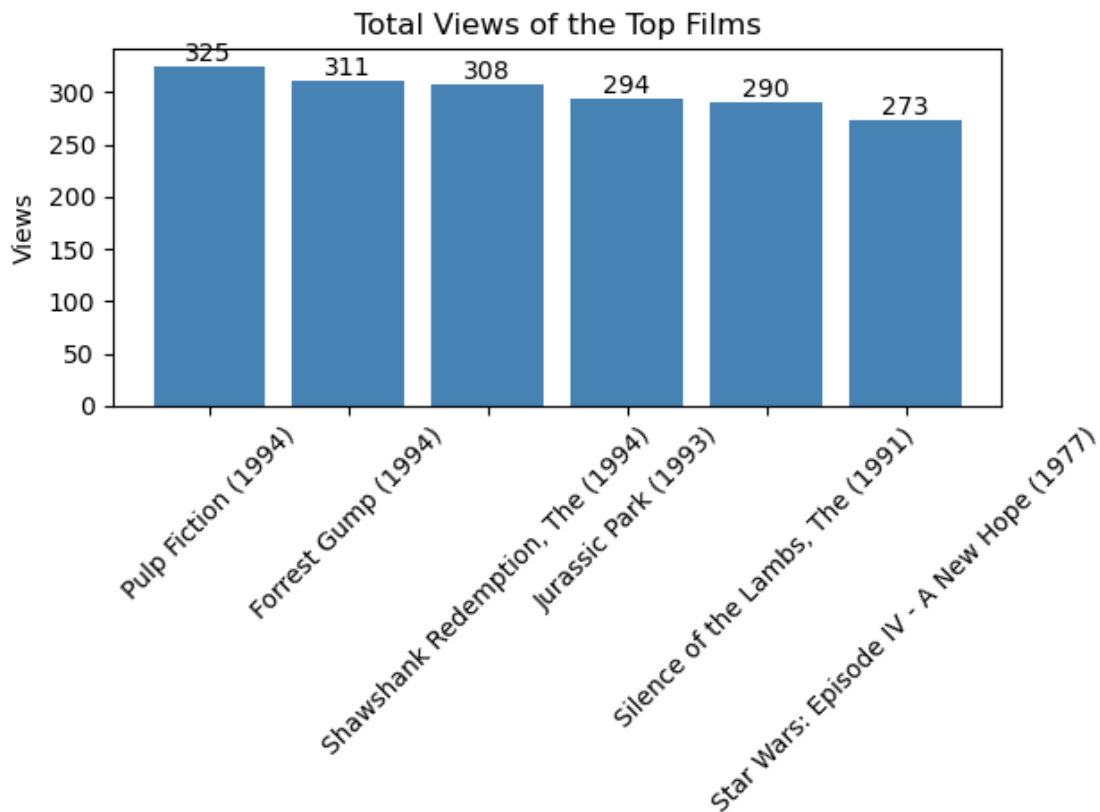
```
[12]: # Step 6: Most Viewed Movies Visualization
# Count how many ratings each movie received (i.e. number of views)
movie_views = rating_matrix.count(axis=0)
table_views = pd.DataFrame({'movieId': movie_views.index, 'views': movie_views.
    ↪ values})
table_views = table_views.sort_values(by='views', ascending=False)

# Merge with movie_data to get movie titles
table_views = table_views.merge(movie_data[['movieId', 'title']], on='movieId',
    ↪ how='left')
print("\nTop 6 Most Viewed Movies:")
print(table_views.head(6))


# Plot the top 6 most viewed movies
top6 = table_views.head(6)
plt.figure()
plt.bar(top6['title'], top6['views'], color='steelblue')
plt.xticks(rotation=45)
for i, v in enumerate(top6['views']):
    plt.text(i, v, str(v), ha='center', va='bottom')
plt.title("Total Views of the Top Films")
plt.ylabel("Views")
plt.tight_layout()
plt.show()
```

Top 6 Most Viewed Movies:

	movieId	views	title
0	296	325	Pulp Fiction (1994)
1	356	311	Forrest Gump (1994)
2	318	308	Shawshank Redemption, The (1994)
3	480	294	Jurassic Park (1993)
4	593	290	Silence of the Lambs, The (1991)
5	260	273	Star Wars: Episode IV - A New Hope (1977)



```
[13]: # Step 7: Heatmap of a Portion of the Rating Matrix
# Display a heatmap of the first 20 users and 25 movies
plt.figure()
plt.imshow(rating_matrix.iloc[0:20, 0:25].fillna(0), aspect='auto',
           cmap='viridis')
plt.title("Heatmap of the first 20 users and 25 movies")
plt.colorbar()
plt.show()
```



```
[15]: # Step 8: Data Preparation - Filtering, Normalization, and Binarization
# Filter out users and movies with fewer than 50 ratings
filtered_users = rating_matrix.count(axis=1) > 50
filtered_movies = rating_matrix.count(axis=0) > 50
movie_ratings = rating_matrix.loc[filtered_users, filtered_movies]
print("\nFiltered Rating Matrix shape:", movie_ratings.shape)

# (a) For further filtering: select top users/movies by count (98th percentile)
min_movies = movie_ratings.count(axis=1).quantile(0.98)
min_users = movie_ratings.count(axis=0).quantile(0.98)
top_users = movie_ratings[movie_ratings.count(axis=1) > min_movies]
top_movies = movie_ratings.loc[:, movie_ratings.count(axis=0) > min_users]

plt.figure()
plt.imshow(top_users[top_movies.columns].fillna(0), aspect='auto',
           cmap='viridis')
plt.title("Heatmap of the Top Users and Movies")
plt.colorbar()
plt.show()

# (b) Distribution of average ratings per user
```



```

average_ratings = movie_ratings.mean(axis=1)
plt.figure()
plt.hist(average_ratings, bins=30, color='steelblue', edgecolor='red')
plt.title("Distribution of the Average Rating per User")
plt.xlabel("Average Rating")
plt.ylabel("Frequency")
plt.show()

# (c) Data Normalization - subtract the mean rating of each user
normalized_ratings = movie_ratings.sub(movie_ratings.mean(axis=1), axis=0)
print("\nNumber of users with non-zero mean (should be near zero):",
      (normalized_ratings.mean(axis=1).abs() > 1e-5).sum())

plt.figure()
plt.imshow(normalized_ratings.loc[normalized_ratings.count(axis=1) > min_movies,
                                 normalized_ratings.count(axis=0) > min_users].
           ↪fillna(0),
           aspect='auto', cmap='viridis')
plt.title("Normalized Ratings of the Top Users")
plt.colorbar()
plt.show()

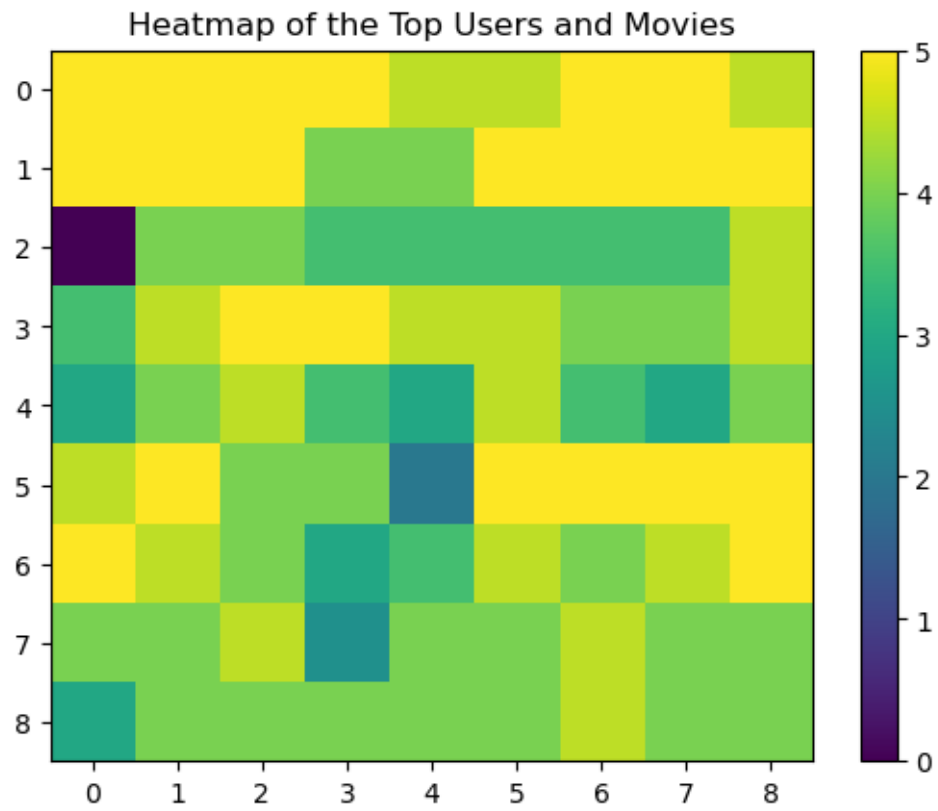
# (d) Data Binarization - convert ratings to 1 if rating >= 3, else 0
binary_min_movies = movie_ratings.count(axis=1).quantile(0.95)
binary_min_users = movie_ratings.count(axis=0).quantile(0.95)
goodRatedFilms = (movie_ratings >= 3).astype(int)


topBinaryUsers = goodRatedFilms[goodRatedFilms.count(axis=1) > ↵
                               ↪binary_min_movies]
topBinaryMovies = goodRatedFilms.loc[:, goodRatedFilms.count(axis=0) > ↵
                                   ↪binary_min_users]

plt.figure()
plt.imshow(topBinaryUsers[topBinaryMovies.columns], aspect='auto', ↵
           ↪cmap='viridis')
plt.title("Heatmap of the Top Users and Movies (Binarized)")
plt.colorbar()
plt.show()


```


Filtered Rating Matrix shape: (420, 447)





Number of users with non-zero mean (should be near zero): 0





```
[16]: # Step 9: Collaborative Filtering - Splitting Data into Training and Testing
      ↪ Sets
      # Randomly split the filtered movie ratings into 80% training and 20% testing
      msk = np.random.rand(len(movie_ratings)) < 0.8
      training_data = movie_ratings[msk]
      testing_data = movie_ratings[~msk]
      print("\nTraining data shape:", training_data.shape)
      print("Testing data shape:", testing_data.shape)
```

Training data shape: (331, 447)

Testing data shape: (89, 447)

```
[17]: # Step 10: Building an Item-Based Collaborative Filtering Model
      # Compute item-item similarity matrix using the training data (fill missing
      ↪ values with 0)
      item_similarity = cosine_similarity(training_data.fillna(0).T)
      item_similarity_df = pd.DataFrame(item_similarity,
                                      index=training_data.columns,
                                      columns=training_data.columns)
      print("\nItem similarity matrix shape:", item_similarity_df.shape)
```

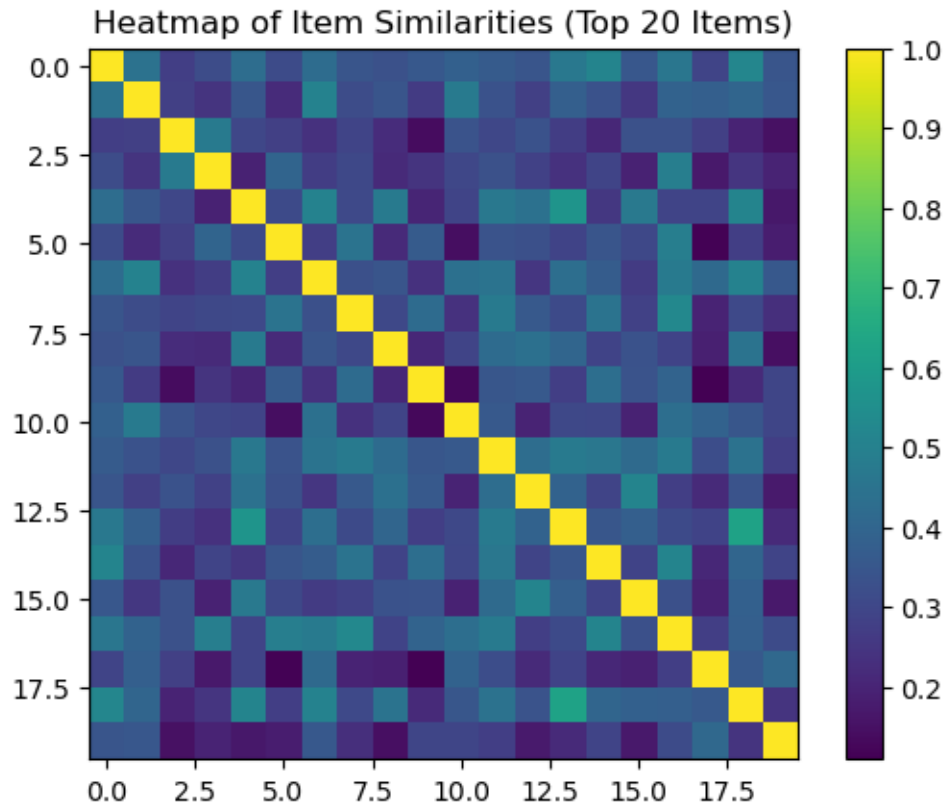
```

# (Optional) Visualize a heatmap of the top 20 items' similarities
top_items = 20
plt.figure()
plt.imshow(item_similarity_df.iloc[:top_items, :top_items],
           ↪interpolation='nearest', cmap='viridis')
plt.title("Heatmap of Item Similarities (Top 20 Items)")
plt.colorbar()
plt.show()

# Define a simple prediction function for a given user using weighted average
def predict_ratings(user_ratings, similarity_df):
    preds = {}
    # For each movie that the user has not rated
    for movie in user_ratings.index:
        if pd.isna(user_ratings[movie]):
            # Consider movies that the user has rated
            rated = user_ratings.notna()
            if rated.sum() == 0:
                preds[movie] = np.nan
            else:
                sims = similarity_df.loc[movie, rated]
                ratings = user_ratings[rated]
                if sims.sum() > 0:
                    preds[movie] = np.dot(sims, ratings) / sims.sum()
                else:
                    preds[movie] = np.nan
    return preds

```

Item similarity matrix shape: (447, 447)



```
[21]: # Step 11: Generate Recommendations for Users in the Test Set
top_n = 10 # number of items to recommend per user
recommendations = {}

# Loop through each user in the test set to generate recommendations
for user_id, user_ratings in testing_data.iterrows():
    preds = predict_ratings(user_ratings, item_similarity_df)
    # Sort the predictions in descending order (ignoring movies with NaN
    ↪ predictions)
    sorted_preds = sorted(preds.items(), key=lambda x: x[1] if not pd.
    ↪ isna(x[1]) else -np.inf, reverse=True)
    recommendations[user_id] = sorted_preds[:top_n]

# Example: Show recommendations for one test user
example_user = list(recommendations.keys())[1]
recommended_movie_ids = [movie for movie, score in
    ↪ recommendations[example_user]]
# Retrieve movie titles from movie_data
```

```
recommended_titles = movie_data[movie_data['movieId'].
    ↪isin(recommended_movie_ids)]['title'].tolist()
print("\nRecommendations for user", example_user, ":")
for i, title in enumerate(recommended_titles, start=1):
    print(f"{i}. {title}")
```

Recommendations for user 16 :

1. Heat (1995)
2. Craft, The (1996)
3. Long Kiss Goodnight, The (1996)
4. Boot, Das (Boat, The) (1981)
5. Star Trek: First Contact (1996)
6. Star Trek IV: The Voyage Home (1986)
7. Sneakers (1992)
8. Last of the Mohicans, The (1992)
9. Saint, The (1997)
10. Payback (1999)

[]: